


Hierarchical methods for Reinforcement Learning

Evgenii Nikishin

30.03.2018

Reinforcement Learning recap

Markov Decision Process (MDP):



- Agent actions $A_t \in \mathcal{A}$
- Environment states $S_t \in \mathcal{S}$
- Reward $R_t \in \mathbb{R}$
- Agent policy $A_t \sim \pi(a|s)$
- State transitions $S_{t+1}, R_{t+1} \sim p(s', r|s, a)$

Optimal policy maximizes the expected discounted return:

$$\pi^* = \arg \max_{\pi} \mathbb{E}_{\pi} [R_1 + \gamma R_2 + \gamma^2 R_3 + \dots] = \arg \max_{\pi} \mathbb{E}_{\pi} \left[\sum_{k=0}^{\infty} \gamma^k R_{k+1} \right]$$

Standard approaches

Value functions:

- $V_{\pi}(s) = \mathbb{E}_{\pi} \left[\sum_{k=0}^{\infty} \gamma^k R_{t+k+1} | S_t = s \right]$
- $Q_{\pi}(s, a) = \mathbb{E}_{\pi} \left[\sum_{k=0}^{\infty} \gamma^k R_{t+k+1} | S_t = s, A_t = a \right]$

Standard approaches

Value functions:

- $V_{\pi}(s) = \mathbb{E}_{\pi} \left[\sum_{k=0}^{\infty} \gamma^k R_{t+k+1} | S_t = s \right]$
- $Q_{\pi}(s, a) = \mathbb{E}_{\pi} \left[\sum_{k=0}^{\infty} \gamma^k R_{t+k+1} | S_t = s, A_t = a \right]$

Properties of value functions:

- $V_{\pi}(s) = \sum_a \pi(a|s) Q_{\pi}(s, a)$
- $Q_{\pi}(s, a) = \sum_{s', r} p(s', r|s, a) [r + \gamma V_{\pi}(s')]$

Standard approaches

Value functions:

- $V_{\pi}(s) = \mathbb{E}_{\pi} \left[\sum_{k=0}^{\infty} \gamma^k R_{t+k+1} | S_t = s \right]$
- $Q_{\pi}(s, a) = \mathbb{E}_{\pi} \left[\sum_{k=0}^{\infty} \gamma^k R_{t+k+1} | S_t = s, A_t = a \right]$

Properties of value functions:

- $V_{\pi}(s) = \sum_a \pi(a|s) Q_{\pi}(s, a)$
- $Q_{\pi}(s, a) = \sum_{s', r} p(s', r|s, a) [r + \gamma V_{\pi}(s')]$
- $V_*(s) = V_{\pi^*}(s) = \max_a \sum_{s', r} p(s', r|s, a) [r + \gamma V_*(s')]$
- $Q_*(s, a) = Q_{\pi^*}(s, a) = \sum_{s', r} p(s', r|s, a) [r + \gamma \max_{a'} Q_*(s', a')]$

Standard approaches

Value functions:

- $V_{\pi}(s) = \mathbb{E}_{\pi} \left[\sum_{k=0}^{\infty} \gamma^k R_{t+k+1} | S_t = s \right]$
- $Q_{\pi}(s, a) = \mathbb{E}_{\pi} \left[\sum_{k=0}^{\infty} \gamma^k R_{t+k+1} | S_t = s, A_t = a \right]$

Properties of value functions:

- $V_{\pi}(s) = \sum_a \pi(a|s) Q_{\pi}(s, a)$
- $Q_{\pi}(s, a) = \sum_{s', r} p(s', r|s, a) [r + \gamma V_{\pi}(s')]$
- $V_*(s) = V_{\pi^*}(s) = \max_a \sum_{s', r} p(s', r|s, a) [r + \gamma V_*(s')]$
- $Q_*(s, a) = Q_{\pi^*}(s, a) = \sum_{s', r} p(s', r|s, a) [r + \gamma \max_{a'} Q_*(s', a')]$

Learn optimal Q-function and act greedily (e.g. Q-learning:

$$Q(s, a) \leftarrow Q(s, a) + \alpha (r(s, a) + \max_{a'} Q(s', a') - Q(s, a))$$

Standard approaches

Parametrize policy $\pi(a|s) = \pi_\theta(a|s)$ within a family of differentiable functions and update w.r.t. its gradient:

$$\nabla_\theta \mathbb{E}_{\pi_\theta} \left[\sum_{k=0}^{\infty} \gamma^k R_{k+1} \right] = \mathbb{E}_{\pi_\theta} [\nabla_\theta \log \pi_\theta(A_t|S_t) Q_{\pi_\theta}(S_t, A_t)]$$

Standard approaches

Parametrize policy $\pi(a|s) = \pi_\theta(a|s)$ within a family of differentiable functions and update w.r.t. its gradient:

$$\nabla_\theta \mathbb{E}_{\pi_\theta} \left[\sum_{k=0}^{\infty} \gamma^k R_{k+1} \right] = \mathbb{E}_{\pi_\theta} [\nabla_\theta \log \pi_\theta(A_t|S_t) Q_{\pi_\theta}(S_t, A_t)]$$

Q_{π_θ} is still unknown:

- use MC estimates
- parametrize and learn (Actor-Critic algorithm)

Options framework in RL

Option ω is a tuple $\langle \mathcal{I}_\omega, \pi_\omega, \beta_\omega \rangle$, where

- $\mathcal{I}_\omega \subseteq \mathcal{S}$ — option initiation set,
- $\pi_\omega(a|s)$ — option policy,
- $\beta_\omega(s)$ — probability of option termination at given state.

When option terminates, control is given to meta-policy $\pi_\Omega(\omega|s)$.

E. g., robot navigation: if there is no obstacle in front (\mathcal{I}_ω), go forward (π_ω) until you get too close to another object (β_ω).

One can think of options as functions in program.

Motivation of options:

- Sample efficiency
- Transfer of knowledge
- Consistent exploration
- Interpretability
- Way to address overfitting in RL

Motivation of options:

- Sample efficiency
- Transfer of knowledge
- Consistent exploration
- Interpretability
- Way to address overfitting in RL ([video](#))

Motivation of options:

- Sample efficiency
- Transfer of knowledge
- Consistent exploration
- Interpretability
- Way to address overfitting in RL ([video](#))

Arguments:

- There exist optimal deterministic policy in MDP
- Introducing options lead to larger number of parameters, which slows down learning

Value functions extension:

$$Q_{\Omega}(s, \omega) = \sum_a \pi_{\omega}(a|s) Q_U(s, \omega, a)$$

Value functions extension:

$$Q_{\Omega}(s, \omega) = \sum_a \pi_{\omega}(a|s) Q_U(s, \omega, a)$$

$$Q_U(s, \omega, a) = \sum_{s', r} p(s', r|s, a) [r(s, a) + \gamma U(\omega, s')]$$

Value functions extension:

$$Q_{\Omega}(s, \omega) = \sum_a \pi_{\omega}(a|s) Q_U(s, \omega, a)$$

$$Q_U(s, \omega, a) = \sum_{s', r} p(s', r|s, a) [r(s, a) + \gamma U(\omega, s')]$$

$$U(\omega, s') = (1 - \beta_{\omega}(s')) Q_{\Omega}(s', \omega) + \beta_{\omega}(s') V_{\Omega}(s')$$

Value functions extension:

$$Q_{\Omega}(s, \omega) = \sum_a \pi_{\omega}(a|s) Q_U(s, \omega, a)$$

$$Q_U(s, \omega, a) = \sum_{s', r} p(s', r|s, a) [r(s, a) + \gamma U(\omega, s')]$$

$$U(\omega, s') = (1 - \beta_{\omega}(s')) Q_{\Omega}(s', \omega) + \beta_{\omega}(s') V_{\Omega}(s')$$

$$V_{\Omega}(s') = \sum_{\omega'} \pi_{\Omega}(\omega'|s') Q_{\Omega}(s', \omega')$$

Value functions extension:

$$Q_{\Omega}(s, \omega) = \sum_a \pi_{\omega}(a|s) Q_U(s, \omega, a)$$

$$Q_U(s, \omega, a) = \sum_{s', r} p(s', r|s, a) [r(s, a) + \gamma U(\omega, s')]$$

$$U(\omega, s') = (1 - \beta_{\omega}(s')) Q_{\Omega}(s', \omega) + \beta_{\omega}(s') V_{\Omega}(s')$$

$$V_{\Omega}(s') = \sum_{\omega'} \pi_{\Omega}(\omega'|s') Q_{\Omega}(s', \omega')$$

$Q_U(s, \omega, a)$ is learnt via Q-learning algorithm.

Main contributions:

Gradient w.r.t. option policy $\pi_{\omega}(a|s)$ params is given by

$$\mathbb{E} [\nabla \log \pi_{\omega}(a|s) Q_U(s, \omega, a)]$$

Natural result: take better primitive actions more often inside the option

Gradient w.r.t. option termination $\beta_{\omega}(s)$ params is given by

$$\mathbb{E} [-\nabla \beta_{\omega}(s) (Q_{\Omega}(s', \omega') - V_{\Omega}(s'))]$$

Also natural: lengthen options that have a large advantage.

Experiments




Figure: 4 rooms domain




Figure: Termination probabilities for the option-critic agent learning with 4 options.

Experiments




Figure: After a 1000 episodes, the goal location in the four-rooms domain is moved randomly.

Stochastic Neural Networks

Consider set of MDPs \mathcal{M} .

Assumption: for each MDP $M \in \mathcal{M}$, state space is decomposed into two components, $\mathcal{S}_{\text{agent}}$, and $\mathcal{S}_{\text{rest}}^M$. Also, MDPs share action space.

E.g. robot facing different tasks.

Stochastic Neural Networks

Consider set of MDPs \mathcal{M} .

Assumption: for each MDP $M \in \mathcal{M}$, state space is decomposed into two components, $\mathcal{S}_{\text{agent}}$, and $\mathcal{S}_{\text{rest}}^M$. Also, MDPs share action space.

E.g. robot facing different tasks.

Scenario:

- 1 One MDP is used to pretrain set of skills.
- 2 For other MDPs, policy on top of this skills is trained.

Skills training

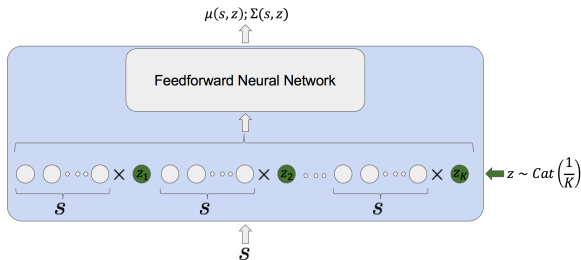


Figure: Skill network architecture

During for each episode separate z is sampled.

Skills training




Figure: Skill network architecture

During for each episode separate z is sampled.
What makes agent to take into account z ?

Mutual Information bonus

$$I(Z; C) = H(Z) - H(Z|C) = \text{const} + \mathbb{E}_{z,c} [\log p(Z = z|C = c)]$$

In order to forbid agent to ignore z , received reward modified:
Mutual Information between z and coordinates at timestep t is added:

$$R_t \leftarrow R_t + \alpha_H \log \hat{p}(Z = z|c_t)$$

In order to estimate $\log \hat{p}(Z = z^n|c_t^n)$, discretization is used: denote $m_c(z)$ visitation counts of how many times each cell c is visited when latent code z is sampled

$$\hat{p}(Z = z|c) = \frac{m_c(z)}{\sum_{z'} m_c(z')}$$

Skills usage




Figure: Hierarchical SNN architecture to solve downstream tasks

Manager operates in terms of skills: z now considered as actions. Once manager selects a skill z , agent is committing to it for a fixed amount of steps \mathcal{T} .

Experimental setup




Figure: MDP 0: locomotion




Figure: MDP 1: Maze 1




Figure: MDP 2: Maze 2








Figure: MDP 3: Food Gather

Experimental setup

- In MDP 0, we are interested in learning diverse set of skills
- Faster learning of the hierarchical architectures in the downstream MDPs:
 - ① CoM reward: single policy with bonus speed reward (authors claim it accelerates learning)
 - ② Multi-policy: independently trained policies + manager network upon
 - ③ SNN
 - ④ SNN + MI bonus

Experiments



Span of learnt skills in MDP 0, $\alpha_H = 0, 0.001, 0.01, 0.1$.

Experiments




Figure: Maze 1




Figure: Maze 2




Figure: Food Gather

References

- ① Between MDPs and semi-MDPs: A framework for temporal abstraction in reinforcement learning (Original options paper)
- ② The Option-Critic Architecture
- ③ Stochastic Neural Networks for Hierarchical Reinforcement Learning