

Variational inference for reinforcement learning in partially observable markov decision processes

Arseny Kuznetsov


Samsung AI center Moscow

April 26, 2019

Overview

- 1 Introduction
 - State space model
 - Sequential Monte-Carlo
 - Variational Sequential Monte-Carlo
- 2 Deep Variational Reinforcement Learning
 - Advantage Actor Critic (A2C)
 - DVRL
- 3 ROBiT
 - Variational inference in MDP
 - ROBiT

State space model



HMM is a tuple (S, T, Ω, O) :

- S – state space
- T – transition function
- Ω – space of observations
- O – observation function

We want to perform filtering, i. e. to estimate

$$p(s_t | o_{1:t}), \quad p(o_{1:t})$$

SMC procedure: Bootstrap filter (Gordon, 1993)

At $t = 1$

① Sample N particles $s_1^i \sim q(s_1)$

② Compute weights $w_1(s_1^i) = \frac{p(s_1^i)p(o_1^i|s_1^i)}{q(s_1^i)}$

$$W_1^i = \frac{w_1(s_1^i)}{\sum_i w_1(s_1^i)}$$

SMC procedure: Bootstrap filter (Gordon, 1993)

At $t = 1$

- 1 Sample N particles $s_1^i \sim q(s_1)$

- 2 Compute weights $w_1(s_1^i) = \frac{p(s_1^i)p(o_1^i|s_1^i)}{q(s_1^i)}$ $W_1^i = \frac{w_1(s_1^i)}{\sum_i w_1(s_1^i)}$

At $t \geq 2$

- 1 Sample ancestor indices $u_{t-1}^i \sim \text{Cat}(W_{t-1}^1, \dots, W_{t-1}^N)$

- 2 Sample N particles $s_t^i \sim q(s_t | s_{1:t-1}^{u_{t-1}^i})$

- 3 Compute weights

$$w_t(s_{1:t}^i) = \frac{p(o_t^i | s_t^i) p(s_t^i | s_{1:t-1}^{u_{t-1}^i})}{q(s_t^i | s_{1:t-1}^{u_{t-1}^i})} \quad W_t^i = \frac{w_t(s_{1:t}^i)}{\sum_i w_t(s_{1:t}^i)}$$

SMC procedure: Bootstrap filter (Gordon, 1993)

At $t = 1$

① Sample N particles $s_1^i \sim q(s_1)$

② Compute weights $w_1(s_1^i) = \frac{p(s_1^i)p(o_1^i|s_1^i)}{q(s_1^i)}$ $W_1^i = \frac{w_1(s_1^i)}{\sum_i w_1(s_1^i)}$

At $t \geq 2$

① Sample ancestor indices $u_{t-1}^i \sim \text{Cat}(W_{t-1}^1, \dots, W_{t-1}^N)$

② Sample N particles $s_t^i \sim q(s_t|s_{1:t-1}^{u_{t-1}^i})$

③ Compute weights

$$w_t(s_{1:t}^i) = \frac{p(o_t^i|s_t^i)p(s_t^i|s_{1:t-1}^{u_{t-1}^i})}{q(s_t^i|s_{1:t-1}^{u_{t-1}^i})} \quad W_t^i = \frac{w_t(s_{1:t}^i)}{\sum_i w_t(s_{1:t}^i)}$$

Estimate normalization constant and target distribution

$$\hat{p}(o_{1:t}) = \prod_{\tau=1}^t \frac{1}{N} \sum_{i=1}^N w_{\tau}(s_{1:\tau}^i) \quad \hat{p}_t(s_{1:t}|o_{1:t}) = \sum_{i=1}^N W_t^i \delta(s_{1:t} - s_{1:t}^i)$$

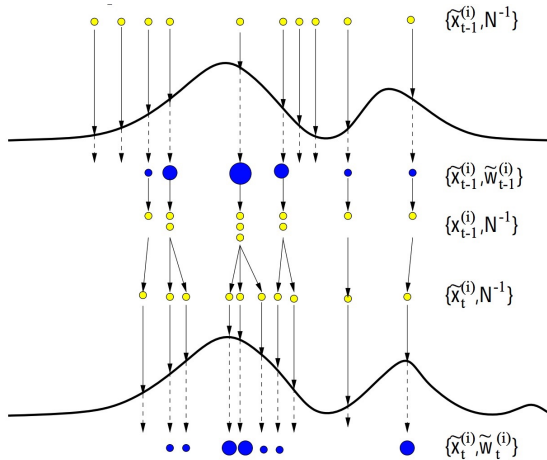


Figure 1. Process of particle filtering

Variational Sequential Monte Carlo

Lets consider learnable model $p_\theta(s_{t+1}|s_t)$, $p_\theta(o_t|s_t)$ and proposal $q_\lambda(s_{t+1}|s_t)$.

$$\begin{aligned} \tilde{\phi}(s_{1:T}^{1:N}, u_{1:T-1}^{1:N}) = \\ \left[\prod_{i=1}^N q(s_1^i) \right] \prod_{t=2}^T \prod_{i=1}^N \left[W_{t-1}^{u_{t-1}^i} q_\lambda(s_t^i | s_{t-1}^{u_{t-1}^i}) \right] \end{aligned} \quad (1)$$


VSMC (Naesseth et al., 2017) shows that

$$\tilde{\mathcal{L}} \triangleq \mathbb{E}_{\tilde{\phi}} \sum_{t=1}^T \log \left(\frac{1}{N} \sum_{i=1}^N w_t^i \right)$$

is a lower bound on $\log p_\theta(o_{1:T})$.

DVRL: Deep Variational Reinforcement Learning for POMDPs

RL in MDP



- A – action space
- $T - p(s_{t+1}|s_t, a_t)$
- Reward distribution $p(r_t|s_t, a_t)$
- $\gamma \in [0, 1]$ – discount factor

Usual denotations:

$$G_t = \sum_{i=0}^{T-t} \gamma^i R_{t+i}$$

$$V^\pi(s) \triangleq \mathbb{E}_\pi [G_t | S_t = s]$$

$$Q^\pi(s, a) \triangleq \mathbb{E}_\pi [G_t | S_t = s, A_t = a]$$

Advantage Actor Critic (A2C)

$$Q^{t,i}(s_{t+i}, a_{t+i}) \triangleq \left(\sum_{j=0}^{n_s-i-1} \gamma^j r_{t+i+j} \right) + \gamma^{n_s-i} V_{\eta}^{-}(s_{t+n_s})$$

Advantage Actor Critic (A2C)

$$Q^{t,i}(s_{t+i}, a_{t+i}) \triangleq \left(\sum_{j=0}^{n_s-i-1} \gamma^j r_{t+i+j} \right) + \gamma^{n_s-i} V_{\eta}^{-}(s_{t+n_s})$$

$$A^{t,i}(s_{t+i}, a_{t+i}) \triangleq Q^{t,i}(s_{t+i}, a_{t+i}) - V_{\eta}(s_{t+i})$$

$$\mathcal{L}_t^A = -\frac{1}{n_e n_s} \sum_{envs} \sum_{i=0}^{n_s-1} \log \pi_{\rho}(a_{t+i}|s_{t+i}) A^{t,i,-}(s_{t+i}, a_{t+i})$$

Advantage Actor Critic (A2C)

$$Q^{t,i}(s_{t+i}, a_{t+i}) \triangleq \left(\sum_{j=0}^{n_s-i-1} \gamma^j r_{t+i+j} \right) + \gamma^{n_s-i} V_{\eta}^{-}(s_{t+n_s})$$

$$A^{t,i}(s_{t+i}, a_{t+i}) \triangleq Q^{t,i}(s_{t+i}, a_{t+i}) - V_{\eta}(s_{t+i})$$

$$\mathcal{L}_t^A = -\frac{1}{n_e n_s} \sum_{envs} \sum_{i=0}^{n_s-1} \log \pi_{\rho}(a_{t+i}|s_{t+i}) A^{t,i,-}(s_{t+i}, a_{t+i})$$

$$\mathcal{L}_t^V = \frac{1}{n_e n_s} \sum_{envs} \sum_{i=0}^{n_s-1} A^{t,i}(s_{t+i}, a_{t+i})^2$$

Advantage Actor Critic (A2C)

$$Q^{t,i}(s_{t+i}, a_{t+i}) \triangleq \left(\sum_{j=0}^{n_s-i-1} \gamma^j r_{t+i+j} \right) + \gamma^{n_s-i} V_{\eta}^{-}(s_{t+n_s})$$


$$A^{t,i}(s_{t+i}, a_{t+i}) \triangleq Q^{t,i}(s_{t+i}, a_{t+i}) - V_{\eta}(s_{t+i})$$

$$\mathcal{L}_t^A = -\frac{1}{n_e n_s} \sum_{envs} \sum_{i=0}^{n_s-1} \log \pi_{\rho}(a_{t+i}|s_{t+i}) A^{t,i,-}(s_{t+i}, a_{t+i})$$

$$\mathcal{L}_t^V = \frac{1}{n_e n_s} \sum_{envs} \sum_{i=0}^{n_s-1} A^{t,i}(s_{t+i}, a_{t+i})^2$$

$$\mathcal{L}_t^H = -\frac{1}{n_e n_s} \sum_{envs} \sum_{i=0}^{n_s-1} H(\pi_{\rho}(\cdot|s_{t+i}))$$

POMDP



- Combination of SSM and MDP.
- We don't have access to underlying states anymore.

Recurrent latent state update

DVRL uses two-part state representation: $s_t^i = (h_t^i, z_t^i)$.
State update basically repeats on VSMC time-step.

- ① $u_{t-1}^k \sim \text{Cat} \left(\frac{w_{t-1}^k}{\sum_{j=1}^K w_{t-1}^j} \right)$
- ② $z_t^k \sim q_\phi(z_t^k | h_{t-1}^{u_{t-1}^k}, a_{t-1}, o_t)$
- ③ $h_t^k = \psi_\theta^{\text{RNN}}(h_{t-1}^{u_{t-1}^k}, z_t^k, a_{t-1}, o_t)$
- ④ $w_t^k = \frac{p_\theta(z_t^k | h_{t-1}^{u_{t-1}^k}, a_{t-1}) p_\theta(o_t | h_{t-1}^{u_{t-1}^k}, z_t^k, a_{t-1})}{q_\phi(z_t^k | h_{t-1}^{u_{t-1}^k}, a_{t-1}, o_t)}$
- ⑤ $\hat{s}_t = \nu_\theta^{\text{RNN}}(\{(h_t^k, z_t^k, w_t^k)\}_{k=1}^K)$


DVRL equals A2C plus PF

$$\mathcal{L}_t^{ELBO} = -\frac{1}{n_e n_s} \sum_{envs} \sum_{i=0}^{n_s-1} \log \left(\frac{1}{K} \sum_{k=1}^K w_{t+i}^k \right)$$

$$\mathcal{L}_t^{DVRL} = \mathcal{L}_t^A + \lambda^H \mathcal{L}_t^H + \lambda^V \mathcal{L}_t^V + \lambda^E \mathcal{L}_t^{ELBO}$$

ROBiT: Joint Reward Optimization and Belief Tracking through Approximate Inference

Variational inference in MDP



O_t – surrogate variable, it indicates "optimality" at moment t .

$$p(O_t = 1|r_t) \propto \exp(r_t)$$

- Trajectory, i.e. all states and actions

$$\mathbf{z} \triangleq \{\mathbf{s}_{1:T+1}, \mathbf{a}_{1:T}\}$$

- Prior over trajectory with respect to prior policy μ

$$p(\mathbf{z}) = p(s_1) \prod_{t=1}^T p(s_{t+1}|s_t, a_t) \mu(a_t|s_t),$$

- Variational posterior over trajectory with respect to policy π

$$q(\mathbf{z}) = p(s_1) \prod_{t=1}^T p(s_{t+1}|s_t, a_t) \pi(a_t|s_t),$$

- Likelihood of optimality given trajectory

$$p(\mathcal{O}_{1:T} = 1 | \mathbf{z}) = \prod_{t=1}^T p(\mathcal{O}_t = 1 | s_t, a_t) \propto \prod_{t=1}^T \exp r(s_t, a_t)$$

- Marginal log likelihood of optimality variable

$$\begin{aligned} \log p(\mathcal{O}_{1:T} = 1) &\geq \mathbb{E}_{q(\mathbf{z})} \log \frac{p(\mathcal{O}_{1:T} = 1 | \mathbf{z}) p(\mathbf{z})}{q(\mathbf{z})} \\ &= \mathbb{E}_{q(\mathbf{z})} \sum_{t=1}^T [r(s_t, a_t) - \text{KL}(\pi(\cdot | s_t) || \mu(\cdot | s_t))] \end{aligned}$$

Soft actor critic

In case of uniform prior lower bound becomes Maximum Entropy objective (Soft Actor Critic, Soft Q-learning):

$$\log p(\mathcal{O}_{1:T} = 1) \geq \mathbb{E}_{q(\mathbf{z})} \sum_{t=1}^T [r(s_t, a_t) + H(\pi(\cdot|s_t))]$$

Soft actor critic:

- Sample efficient, best scores
- Insensitive to hyperparameter change

Why does it work well?

- Tries to find several way to solve a task, multimodal policies
- Finds conservative, safe solutions

ROBiT lower bounds

Lets denote:

$$x_t = [r_t, \mathcal{O}_t, o_{t+1}]$$

$$g_t = [r_{1:t-1}, a_{1:t-1}, o_{1:t}, \mathcal{O}_{1:t-1}]$$

ROBiT lower bounds

Lets denote:

$$x_t = [r_t, \mathcal{O}_t, o_{t+1}]$$

$$g_t = [r_{1:t-1}, a_{1:t-1}, o_{1:t}, O_{1:t-1}]$$

Construct lower bound on marginal loglikelihood:

$$\begin{aligned} L_t &= \log p(x_t, \mathcal{O}_{t+1:T} | g_t) \\ &= \log \mathbb{E}_{s_t} [p(x_t, \mathcal{O}_{t+1:T} | s_t)] \\ &\geq \mathbb{E}_{s_t} [\log p(x_t, \mathcal{O}_{t+1:T} | s_t)] \\ &\geq \mathbb{E}_{s_t} \mathbb{E}_{a_t, s_{t+1} \sim q} \left[\log \frac{p(a_t, x_t, s_{t+1}, \mathcal{O}_{t+1:T} | s_t)}{\pi(a_t | s_t) q(s_{t+1} | s_t, a_t, x_t)} \right] \end{aligned}$$

$$\mathbb{E}_q \left[\underbrace{\log p(\mathcal{O}_{t+1:T} | s_{t+1})}_{\text{value estimate}} + \underbrace{\log p(\mathcal{O}_t | r_t)}_{\text{current reward}} - \underbrace{\text{KL}(\pi | \mu)}_{\text{regularization}} + \underbrace{\log \bar{w}_t}_{\text{VAE on state}} \right]$$

$$\text{where } \bar{w}^t = \frac{p(r_t | s_t, a_t) p(s_{t+1} | s_t, a_t) p(o_{t+1} | s_{t+1}, a_t)}{q(s_{t+1} | s_t, a_t, r_t, o_{t+1})}$$

$$\mathbb{E}_q \left[\underbrace{\log p(\mathcal{O}_{t+1:T}|s_{t+1})}_{\text{value estimate}} + \underbrace{\log p(\mathcal{O}_t|r_t)}_{\text{current reward}} - \underbrace{\text{KL}(\pi|\mu)}_{\text{regularization}} + \underbrace{\log \bar{w}_t}_{\text{VAE on state}} \right]$$

$$\text{where } \bar{w}^t = \frac{p(r_t|s_t, a_t)p(s_{t+1}|s_t, a_t)p(o_{t+1}|s_{t+1}, a_t)}{q(s_{t+1}|s_t, a_t, r_t, o_{t+1})}$$

We approximate future optimality with value function.

$$V_\varphi(s_{t+1}) \approx \log p(\mathcal{O}_{t+1:T}|s_{t+1})$$

$$\pi(a|g_t) = \sum_{k=1}^K W_t^{k,N} \pi_\phi(a|s_t^k).$$

$$V(g_t) = \frac{1}{K} \sum_{k=1}^K V_\varphi(s_t^k)$$

- Naive particle filter implementation would need to have multiple a_t^k particles.
- But we can have only one action in the environment.

SIVI bound

We utilize a variant of SIVI lower bound. Additional state particles are sampled to estimate action probability.

$$\tilde{s}_t^{0:N} \sim p_\theta(s_t|g_t), \quad a_t \sim \pi(\tilde{s}_t^0) \quad (2)$$

$$\pi(a_t|\tilde{s}_t^{0:N}) = \frac{1}{N+1} \sum_{n=0}^N \pi(a_t|\tilde{s}_t^n). \quad (3)$$

$$w_{t+1}^{k,N}(s_t^k, a_t, s_{t+1}^k) = \frac{p_\theta(a_t, x_t, s_{t+1}^k | s_t^k)}{\pi(a_t|\tilde{s}_t^{0:N})q(s_{t+1}^k | s_t^k, a_t, x_t)} \quad (4)$$

$$\mathcal{L}_t^{\text{SI}} = \mathbb{E}_{\tilde{s}_t^{0:N}, s_t^{1:K}, a_t, s_{t+1}^{1:K}} \left[\log \frac{1}{K} \sum_{k=1}^K w_{t+1}^{k,N}(s_t^k, a_t, s_{t+1}^k) \right]$$

where the expectation is taken with respect to

$$\prod_{n=0}^N p_{\theta}(\tilde{s}_t^n | g_t) \pi(a_t | \tilde{s}_t^0) \prod_{k=1}^K p_{\theta}(s_t^k | g_t) q_{\phi}(s_{t+1}^k | s_t^k, a_t, x_t), \quad (5)$$

$\mathcal{L}_t^{\text{SI}}$ is still a lower bound on L_t

$$\mathcal{L}^{\text{ROBiT}} = \mathcal{L}_t^{\text{SI}} + \lambda^V \mathcal{L}_t^V \quad (6)$$

Recap of the method

- 1 Theoretically grounded belief agregation
- 2 Showed that three separate DVRL losses can be viewed as one joint loss
- 3 From practical viewpoint we added reward likelihood into particle filtering

Experiments




Figure 2. Mountain Hike environment

ROBiT-R – described method.

ROBiT-B – μ is not uniform, but exponential average of π .




Figure 3. Comparison of returns on Mountain Hike.

The results for two variations of ROBiT method and DVRL for flickering Atari environments.

Env	ROBiT-B	ROBiT-G	DVRL
Asteroids	1627 ± 47	1598 ± 58	1539 ± 73
BeamRider	3294 ± 123	2039 ± 288	1663 ± 183
Bowling	27.4 ± 1.0	30.4 ± 0.1	29.5 ± 0.2
Centipede	4550 ± 190	3848 ± 149	4240 ± 116
ChopperCommand	9895 ± 536	8313 ± 1086	6602 ± 449
DoubleDunk	-11.2 ± 0.9	-6.3 ± 6.2	-6.0 ± 1.1
Frostbite	292 ± 12	279 ± 12	297 ± 7
IceHockey	-5.1 ± 0.1	-4.3 ± 0.1	-4.9 ± 0.2
MsPacman	2512 ± 458	2238 ± 103	2221 ± 199
Pong	14.7 ± 3.8	11.6 ± 2.2	18.2 ± 2.7

Thank you!

References

- Naesseth, Christian A., et al. "Variational Sequential Monte Carlo." arXiv preprint arXiv:1705.11140 (2017).
- Le, Tuan Anh, et al. "Auto-Encoding Sequential Monte Carlo." arXiv preprint arXiv:1705.10306 (2017).
- Igl, Maximilian, et al. "Deep Variational Reinforcement Learning for POMDPs." arXiv preprint arXiv:1806.02426 (2018)
- Levine, Sergey. "Reinforcement Learning and Control as Probabilistic Inference: Tutorial and Review. " arXiv preprint arXiv:1805.00909 (2018)