


# Quantization of neural networks

Alexander Fritzer

# How to make inference of the neural network?



GPU



NPU

# How to make inference of the neural network?



GPU

float32




NPU

int8

# How numbers are stored in memory?


**int8**



Sign bit:

0 for positive, 1 for negative

**float32**



# The problem: the quality falls significantly



## How to fix it?

# Let's consider a vector

Vector of floats

**W**

float32



Vector of numbers  
from **0** to **255**

How to transfer it?

# Let's consider a vector

## Asymmetric quantization

We need: scale-parameter (  $\Delta$  ) and a zero-point (  $z$  )

## Quantization operation

$$x_{int} = round\left(\frac{x}{\Delta}\right) + z$$

$$x_Q = clamp(0, N_{levels} - 1, x_{int})$$

## Dequantization operation

$$x_{float} = (x_Q - z)\Delta$$

# Let's consider a vector

## Symmetric quantization

We need: scale-parameter ( $\Delta$ )

## Quantization operation

$$x_{int} = round\left(\frac{x}{\Delta}\right)$$

$$x_Q = clamp(-N_{levels}/2, N_{levels}/2 - 1, x_{int}) \quad \text{if signed}$$

$$x_Q = clamp(0, N_{levels} - 1, x_{int}) \quad \text{if un-signed}$$


## Dequantization operation

$$x_{out} = x_Q \Delta$$



# How to choose the parameters?

**It's easy: based on the distribution**



# How do we quantize?

## **Weights**

Based on the minimum and maximum weight

We can do

- 1) Per-layer quantization
- 2) Per-channel (per-neuron) quantization

## **Activations**

**Requires data.**

Based on the minimum and maximum (or some quantiles) of activations.

# Quantized convolution

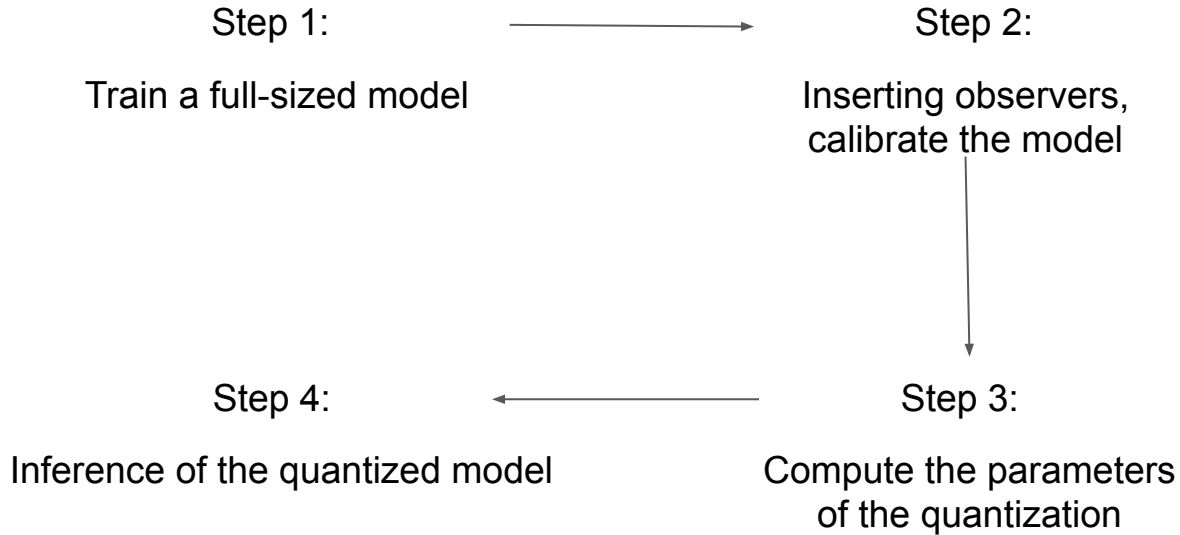
**Naive:**

$$y(k, l, n) = \Delta_w \Delta_x \text{conv}(w_Q(k, l, m; n) - z_w, x_Q(k, l, m) - z_x)$$

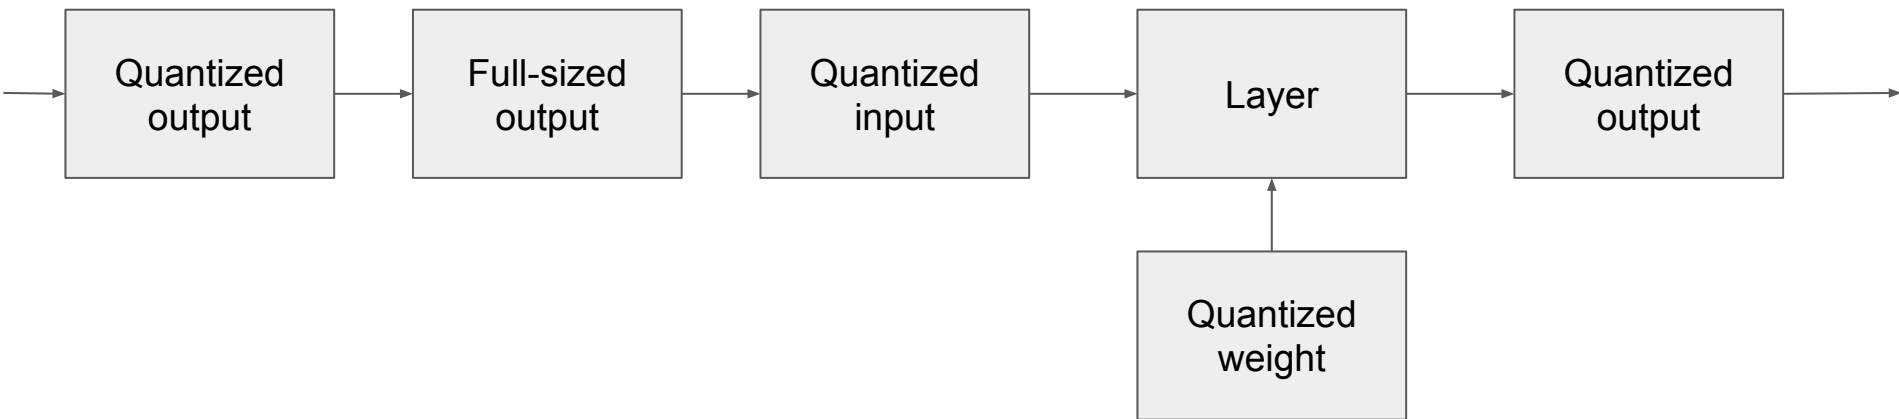
**Advanced:**

$$\begin{aligned} y(k, l, n) = & \text{conv}(w_Q(k, l, m; n), x_Q(k, l, m)) - z_w \sum_{k=0}^{K-1} \sum_{l=0}^{K-1} \sum_{m=0}^{N-1} x_Q(k, l, m) \\ & - z_x \sum_{k=0}^{K-1} \sum_{l=0}^{K-1} \sum_{m=0}^{N-1} w_Q(k, l, m; n) + z_x z_w \end{aligned}$$

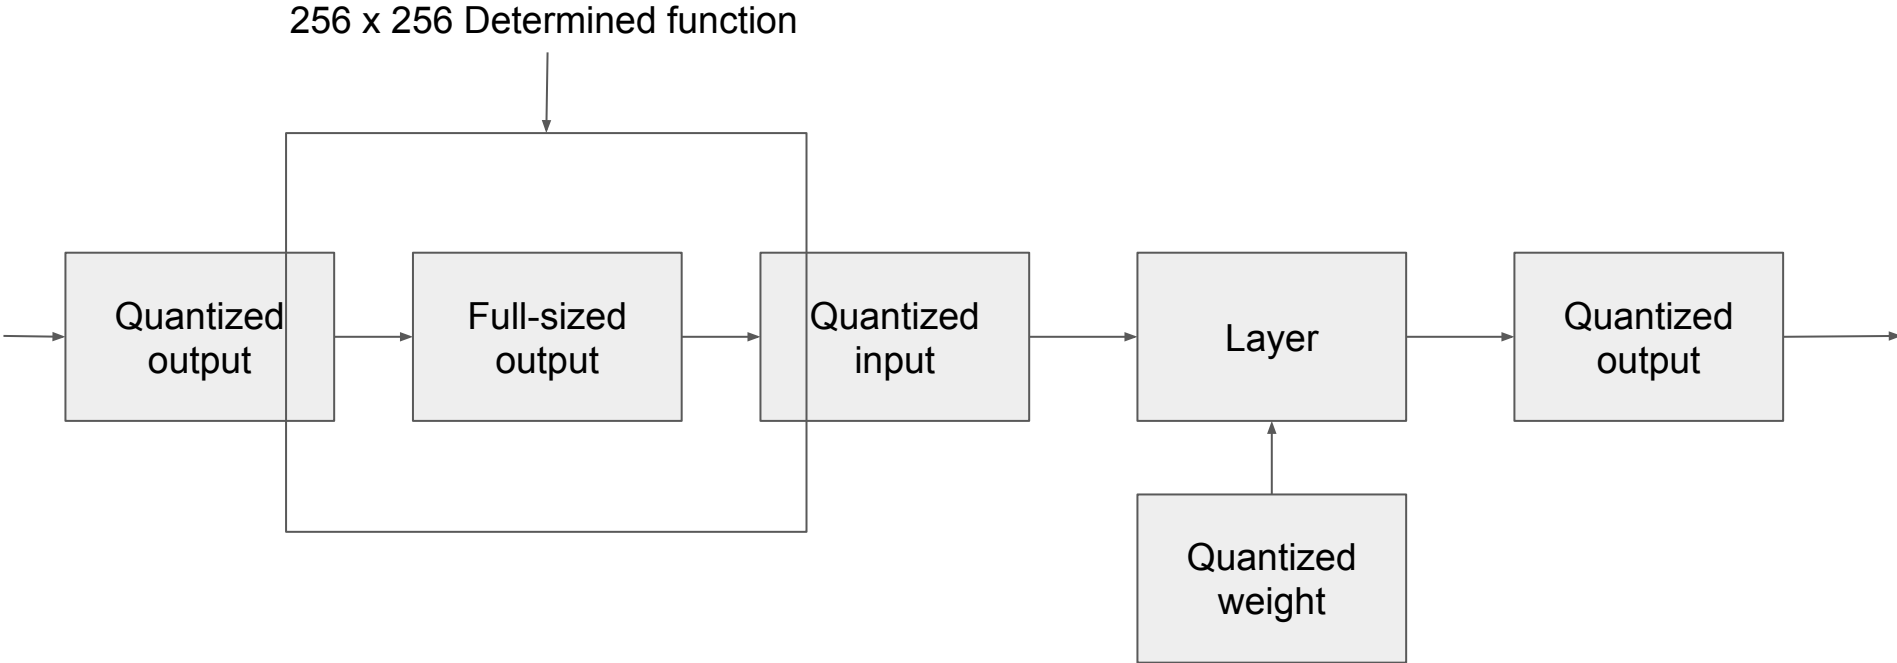
# Post-training quantization pipeline



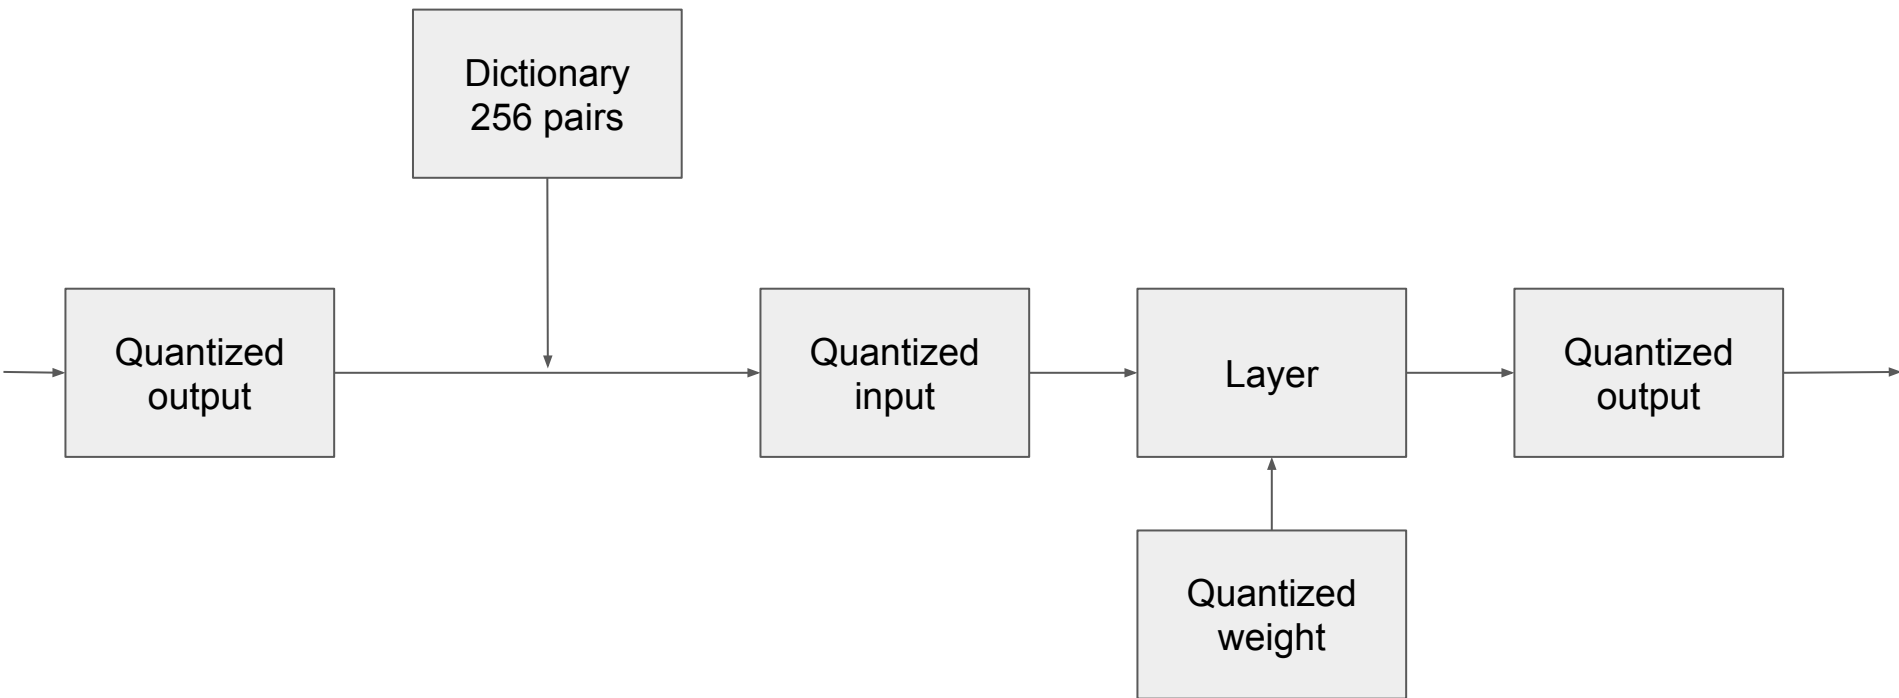
# Quantized inference: theory and practice



# Quantized inference: theory and practice



# Quantized inference: theory and practice




# Results

Network	Asymmetric, per-layer	Symmetric , per-channel	Asymmetric, per-channel	Floating Point
Mobilenetv1_1_224	0.001	0.591	0.704	0.709
Mobilenetv2_1_224	0.001	0.698	0.698	0.719
NasnetMobile	0.722	0.721	0.74	0.74
Mobilenetv2_1.4_224	0.004	0.74	0.74	0.749
Inceptionv3	0.78	0.78	0.78	0.78
Resnet_v1_50	0.75	0.751	0.752	0.752
Resnet_v2_50	0.75	0.75	0.75	0.756
Resnet_v1_152	0.766	0.763	0.762	0.768
Resnet_v2_152	0.761	0.76	0.77	0.778



# Weights distribution



# Bias correction

$$\tilde{\mathbf{y}} = \widetilde{\mathbf{W}}\mathbf{x} \quad - \text{noised output}$$

$$\tilde{\mathbf{y}} = \mathbf{y} + \boldsymbol{\epsilon}\mathbf{x}, \text{ where } \boldsymbol{\epsilon} = \widetilde{\mathbf{W}} - \mathbf{W}$$

$\mathbb{E}[\boldsymbol{\epsilon}\mathbf{x}]_i \neq 0 \longrightarrow$  The expectation of the output will be different

**Solution:** Compute  $\mathbb{E}[\boldsymbol{\epsilon}\mathbf{x}]$  empirically using data and subtract it from bias

# Results

Model	MobileNetV2 SSD-lite	DeeplabV3+ (MobileNetV2 backend)
Original model	10.63	41.40
DFQ (ours)	<b>67.91</b>	<b>72.33</b>
Per-channel quantization	67.52	71.44
Original model ( <b>FP32</b> )	68.47	72.94

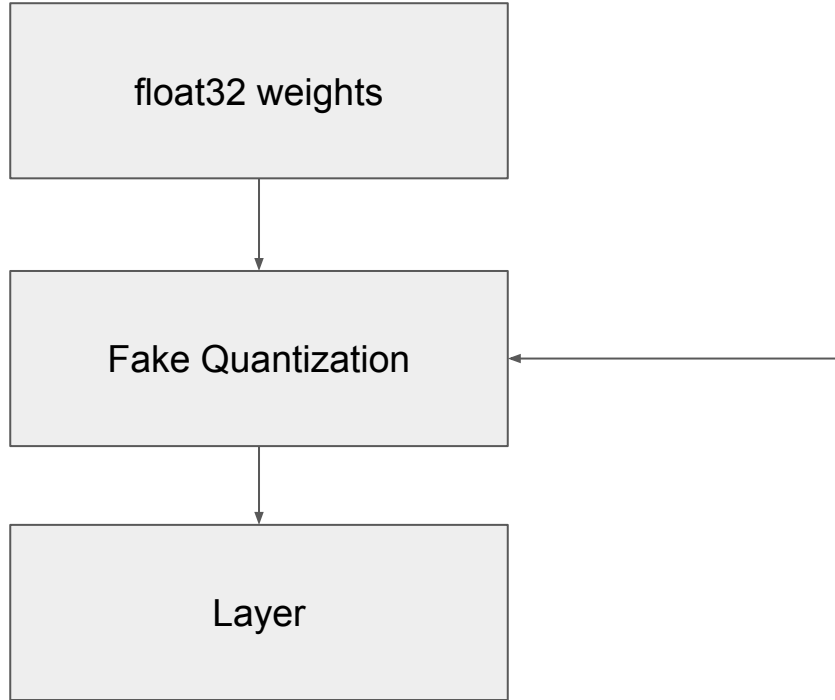
# Can we do better?

- Can we make the gap between the quantized model and full-sized model even smaller?
- Some models do not perform properly after the quantization
- Can we **train quantized model on our computer?**


# “Quantization-Aware Training”

- How the model is stored?
- How do we do forward pass?
- How do we do backpropagation?
- How we choose the parameters of the quantization?

# How the model is stored



Quantization function




# How do we do forward pass (with Pytorch)

	Works on GPU	8-bit computations	8-bit weights and activations
<code>nn.quantized.modules.linear.Linear</code> ( <code>nn.quantized.modules.conv.Conv2d</code> )			
<code>nn.qat.modules.linear.Linear</code> ( <code>nn.qat.modules.conv.Conv2d</code> )			


# How do we do backpropagation

Quantization function (forward)



**Derivative = 0**

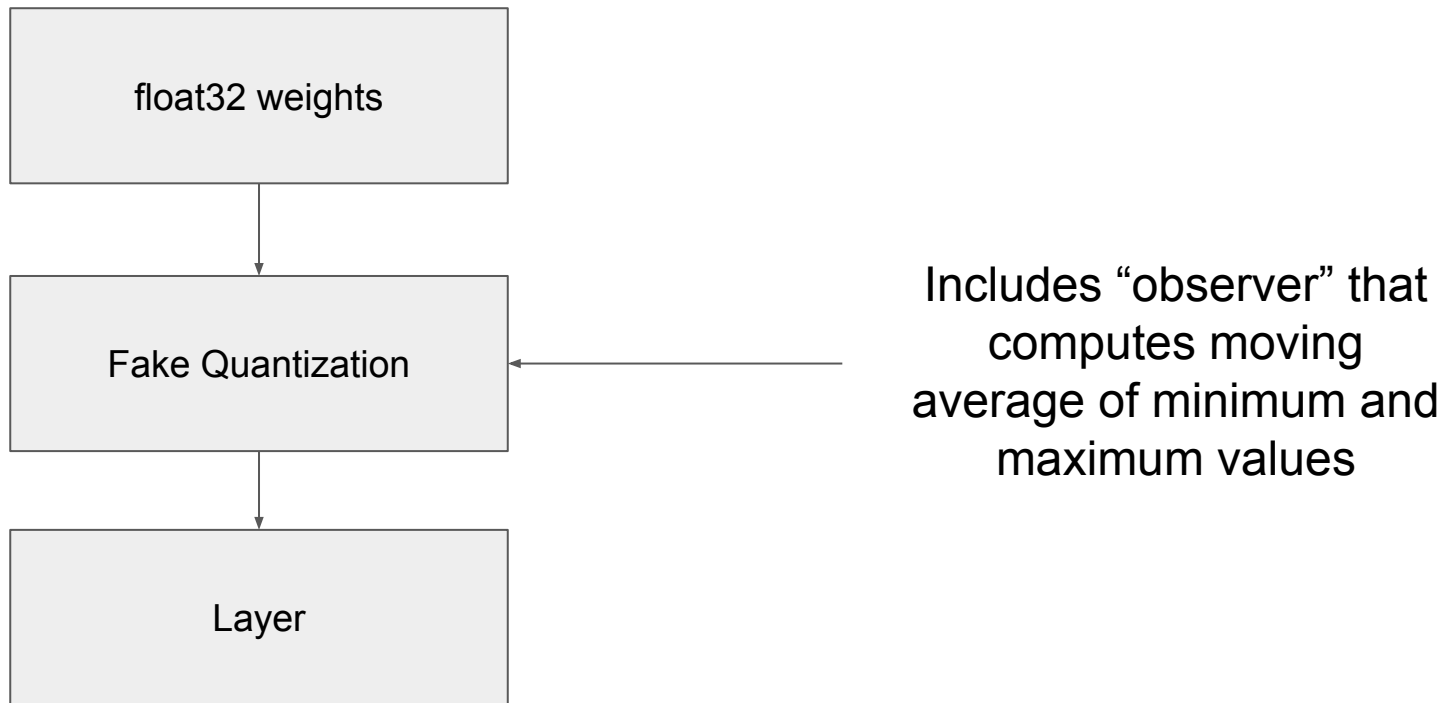
Quantization function (backward)



**Derivative = const**



# How we choose the parameters of the quantization



# Comparison of the inference time

Network	Inference Platform		
	Floating point(CPU)	Fixed point (CPU)	Fixed point (HVX, NN-API)
Mobilenet_v1_1_224	155	68	16
Mobilenet_v2_1_224	105	63	15.5
Mobilenet_v1_1_224_SSD	312	152	
Inception_v3	1391	536	
Resnet_v1_50	874	440	
Resnet_v2_50	1667	1145	
Resnet_v1_152	2581	1274	
Resnet_v2_152	4885	3240	

# References

- Raghuraman Krishnamoorthi. Quantizing deep convolutional networks for efficient inference: A whitepaper. arXiv preprint arXiv:1806.08342, Jun 2018.
- Markus Nagel, Mart van Baalen, Tijmen Blankevoort, and Max Welling. Data-free quantization through weight equalization and bias correction. arXiv preprint arXiv:1906.04721, 2019
- B. Jacob, S. Kligys, B. Chen, M. Zhu, M. Tang, A. Howard, H. Adam, and D. Kalenichenko, “Quantization and Training of Neural Networks for Efficient Integer-Arithmetic-Only Inference,” Dec. 2017.
- <https://pytorch.org/docs/stable/quantization.html>