

Intrinsic Motivation and Automatic Curricula via Asymmetric Self-Play

D. Sharafyan

<https://arxiv.org/pdf/1703.05407.pdf>

Table of contents

1 Asymmetric self-play

2 Policy gradient

3 Experiments

Motivation

Cons of supervised approaches to reinforcement learning

- requires a huge number of episodes to learn
- agent learns the environment at the same time as it tries to understand which trajectories lead to rewards

The idea is to use unsupervised training for an agent, so it can learn the environment without any external reward

Self-play

Agent has two separate "minds": Alice and Bob
During self-play episodes

- Alice's job is to propose a task to Bob
- Bob's job is to complete given task

When presented with a target task episode, Bob is used to perform it

Self-play

Given approach restricted to two classes of environment:

- reversible
- ones that can be reset to their initial state

Self-play

Alice starts at state s_0 and proposes task by doing it, i.e. executing a sequence of actions that takes the agent to a state s_t . Then she outputs a STOP action and hands control over to Bob.

Now there's two options

- if environment is reversible then Bob's task is to return from s_t to s_0
- else Alice's STOP action reinitializes environment and Bob's task is to repeat Alice's task, i.e. reach s_t from s_0

Rewards for Bob and Alice

Let t_A be the time until Alice performs STOP action and t_B is the time taken by Bob to complete his task. The total length of episode is limited to t_{Max} , so if Bob fails to complete his task in time we set $t_B = t_{Max} - t_A$. Bob's reward is

$$R_B = -\gamma t_B$$

And Alice's reward is

$$R_A = \gamma \max(0, t_B - t_A)$$

So Alice is trying to find the simplest task that Bob cannot complete

Parametrizing Alice and Bob's actions

In Alice's case the function will be

$$a_A = \pi_A(s_t, s_0)$$

where s_0 is the observation of the initial state of the environment and s_t is the observation of current state. In Bob's case the function will be

$$a_B = \pi_B(s_t, s^*)$$

where s^* is the target state that Bob has to reach

In all experiments authors use policy gradient with a baseline for optimizing policies.

Table of contents

1 Asymmetric self-play

2 Policy gradient

3 Experiments

Definitions

$$Q(s, a) = \mathbb{E}_{s_{t+1}, a_{t+1}, \dots} \left[\sum_{l=0}^{\infty} \gamma^l r_{t+l} \right]$$

$$V_{\pi}(s_t) = \mathbb{E}_{a_t, s_{t+1}, \dots} \left[\sum_{l=0}^{\infty} \gamma^l r_{t+l} \right]$$

$$A_{\pi}(s, a) = Q_{\pi}(s, a) - V_{\pi}(s)$$

where $a_t \sim \pi(a_t|s_t)$, $s_{t+1} \sim P(s_{t+1}|s_t, a_t)$

Policy objective function

- Given policy $\pi_\theta(a|s)$ with parameters θ , find best θ
- So we need somehow measure the quality of policy $\pi_\theta(a|s)$
- We can use expected discounted reward

$$J(\theta) = \mathbb{E}_{s \sim p(s), a \sim \pi_\theta(s|a)} Q(s, a) = \int_s p(s) \int_a \pi_\theta(a|s) Q(s, a) da ds$$

$$\begin{aligned} \nabla J(\theta) &= \int_s p(s) \int_a \nabla \pi_\theta(a|s) Q(s, a) da ds \\ &= \int_s p(s) \int_a \pi_\theta(a|s) \nabla \log \pi_\theta(a|s) Q(s, a) da ds \\ &= \mathbb{E}_{s \sim p(s), a \sim \pi_\theta(s|a)} \nabla \log \pi_\theta(a|s) Q(s, a) \end{aligned}$$

REINFORCE

Initialize NN weights θ_0

for $i = 1$ to m **do**

 Sample N sessions Z under current $\pi_\theta(a|s)$

 Evaluate policy gradient

$$\nabla J = \frac{1}{N} \sum_{i=0}^N \sum_{s,a \in Z_i} \nabla \log \pi_\theta(a|s) Q(s, a)$$

Update weights

$$\theta_{i+1} = \theta_i + \alpha \nabla J$$

Policy gradient with baseline

$$\Delta\theta = \sum_{t=1}^T \left[\frac{\partial \log \pi_{\theta}(a_t | s_t)}{\partial \theta} \left(\sum_{i=t}^T r_i - b(s_t, \theta) \right) - \lambda \frac{\partial}{\partial \theta} \left(\sum_{i=t}^T r_i - b(s_t, \theta) \right)^2 \right]$$

where $b(s_t, \theta)$ is a baseline function, computed via an extra head on the network producing the action probabilities.

Table of contents


1 Asymmetric self-play

2 Policy gradient


3 Experiments

Long hallway

The environment consists of M states $\{s_1, s_2, \dots, s_M\}$ arranged as a chain. There is 3 actions: "left", "right", "stop". The goal is to reach target state and execute stop action. Here $M = 25$ and $t_{Max} = 30$. Exploration gives bonus reward $\frac{\alpha}{N_s}$ where N_s is the number of times the agent has been in state s




Mazebase: Light key




The maze contains a light switch, a key and a wall with a door. An agent can open or close the door by toggling the key switch, and turn on or off light with the light switch. When the light is off, the agent can only see the (glowing) light switch.

Mazebase results



Mazebase results



StarCraft: Training Marines

The target task is to build Marine units



StarCraft results

Count-based exploration is similar to the hallway experiment.

