

# Natural Language Processing 2

Pavel Yurlov

Higher School of Economics  
Faculty of Computer Science

December 14, 2018

# Outline

- 1 Convolutional networks
  - Summary
  - Interpretation
- 2 CNNs vs. RNNs
  - TCN
- 3 Graph2Seq Encoder
  - Beyond Seq2Seq
  - Model description
  - Comparison

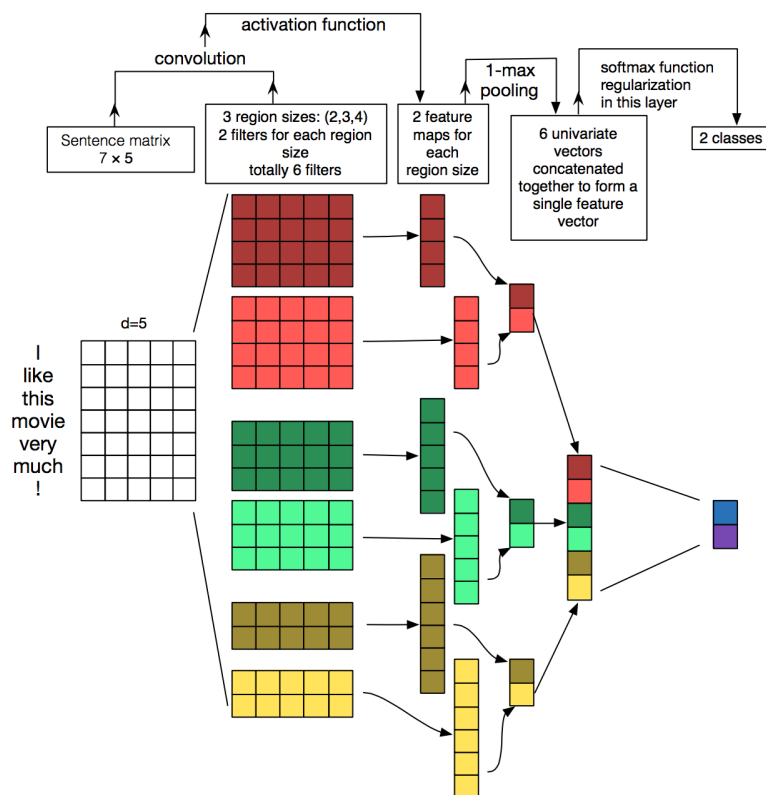
# Convolutional networks

# Convolutional neural networks in NLP

Instead of pixels, the input to most NLP tasks is sentences and documents represented as matrices (e. g., each row is a vector which represents a word).

Thus, NLP filters slide over full rows of the matrix, and their width is the same as the matrix's.

# Scheme



Source: [www.wildml.com/2015/11/understanding-convolutional-neural-networks-for-nlp/](http://www.wildml.com/2015/11/understanding-convolutional-neural-networks-for-nlp/)

# Pooling

- Fixed output size provided by pooling allows us to use variable input or filter sizes
- Detection of specific features without their location

# Channels

Different representations of the same data

- Visual recognition CNNs: RGB channels
- CNNs in NLP: different word embeddings, or different languages, or sentences phrased differently

# Interpretability

Interpretability:

- ① Model interpretability — structured explanation of how the model learns
- ② Prediction interpretability — explanation of how the model arrived at its prediction

Compared to prediction interpretation, neural NLP model interpretation remains under-explored.



# Common explanations

- 1-dimensional convolving filters detect ngrams
- max-pooling extracts the relevant ngrams
- the rest of the network classifies the text based on this information

# Model for text classification

input words  $w_1, \dots, w_n \rightarrow$  vectors  $w_1, \dots, w_n \in \mathbb{R}^d$

$$u_i = [w_i, \dots, w_{i+l-1}]$$

$$F_{ij} = \langle u_i, f_j \rangle$$

$$p_j = \text{ReLU}(\max_i \{F_{ij}\})$$

$$o = \text{softmax}(Wp)$$

Datasets for Sentiment Analysis

# Informative vs. uninformative ngrams

- Assumption (validated): there is a threshold for each filter, items below which are uninformative
- Deliberate vs. accidental ngrams
- Correlation between  $p_j$  and the predicted label for vector  $p$ :  
Filter  $f_j$  contributes to class  $c_j = \arg \max_k W_{kj}$  (in linear case)  
Correlation label — compare  $c_j$  and the final decision by the network

Classifier over a set of texts: pooled vectors  $p^i$  and network predictions  $c^i$

$\forall i : \text{dataset}(X, Y)_j = \{(p_j^i, c^i = c_j) | j < m \& i < D\}$

threshold  $t_j$   $\text{purity}(f, t) = \frac{|\{(x, y) \in (X, Y)_j | x \geq t \& y = \text{true}\}|}{|\{(x, y) \in (X, Y)_j | x \geq t\}|}$

# Purity vs. test accuracy




Figure 1: Evaluation results for identifying important ngrams on the MR model.

Source: [arxiv.org/abs/1809.08037](https://arxiv.org/abs/1809.08037)

# What is captured by a filter?

Intuition (challenged): each feature is homogeneous.

- Slot activation vectors:  $(\langle w_1, f(1) \rangle, \dots, \langle w_l, f(l) \rangle)$

$$\text{Ngram score } \langle u, f \rangle = \sum_{i=0}^{l-1} \langle w_i, f_{id:i(d+1)} \rangle$$

Slot  $i$  of the filter weight  $f(i) = f_{id:i(d+1)}$

- On observed ngrams, the filters achieve maximum values only on some slots.

Two hypotheses:

- Each filter captures multiple semantic classes of ngrams (clustering deliberate ngrams according to their slot activation patterns)
- A slot is used not to detect word existence, but lack thereof

# Common explanations, refined

## ① Filters

- Heterogeneous (a single filter may detect different families of ngrams)
- Detect negative items in ngrams

## ② Max-pooling

- Induces a threshold behaviour, with values below a given threshold ignored

# Practical use

- ① Model interpretability:  
“Visualise” each filter:
  - Class to which its strong signals correspond to
  - Threshold value, purity and coverage percentage
  - List of semantic patterns, each item corresponds to a slot-activations cluster, each cluster to top-k ngrams, each ngram to its total activation, slot-activating vector, list of bottom-k negative ngrams with their activations and slot activations
- ② Prediction interpretability is improved by focusing on informative ngrams and taking into account the negative ones

# CNNs vs. RNNs



# Sequence Modelling

$$f : \mathcal{X}^{T+1} \rightarrow \mathcal{Y}^{T+1}$$

$$\hat{y}_0, \dots, \hat{y}_T = f(x_0, \dots, x_T)$$

Causal constraint:  $y_t$  depends only on  $x_0, \dots, x_t$

$$L(y_0, \dots, y_T, f(x_0, \dots, x_T)) \rightarrow \min_f$$

This formalism encompasses auto-regressive prediction, but does not capture sequence-to-sequence prediction in general, since the entire input sequence can be used to predict each output.

# TCN

Temporal convolutional network:

- ① can take an input of any length and map it to an output sequence of the same length
- ② causal convolutions, no information “leakage” from future to past

Also, a combination of very deep networks and dilated convolutions is used to build very long effective history sizes.


# 1D FCN

TCN = 1D FCN + causal convolutions

1D fully-convolutional (all layers are convolutional, none are fully-connected) network:

- each hidden layer is of the same length as the input layer
- zero padding of length (filter size - 1)


# Causal convolutions



Source: [arxiv.org/abs/1803.01271](https://arxiv.org/abs/1803.01271)

# Dilated convolutions

$$F(s) = (x *_d f)(s) = \sum_{i=0}^{k-1} f(i) \cdot x_{s-d \cdot i}$$




Source: [arxiv.org/abs/1803.01271](https://arxiv.org/abs/1803.01271)

# Residual connections

$$o = \text{Activation}(x + \mathcal{F}(x))$$

Allows networks to learn modifications on the identity mapping which has repeatedly been shown beneficial.



Source: [arxiv.org/abs/1803.01271](https://arxiv.org/abs/1803.01271)

# Pros and cons

## Advantages:

- parallelism (contra RNNs)
- flexible receptive field size
- stable gradients (contra ordinary RNNs)
- low memory requirement for training (no need to store the partial results, the filters are shared across a layer)
- variable length inputs

## Disadvantages:

- data storage during evaluation (need to take in the raw sequence up to the effective history length; RNNs only need to maintain a hidden state)
- potential parameter change for domain transfer (different requirements on the amount of history needed)

# Tasks

- ① Adding problem\*
- ② Sequential MNIST and P-MNIST\*
- ③ Copy memory\*
- ④ JSB Chorales and Nottingham
- ⑤ PennTreebank
- ⑥ Wikitext-103
- ⑦ LAMBADA
- ⑧ text8

\*: synthetic stress tests




# Results

Sequence Modeling Task	Model Size ( $\approx$ )	Models			
		LSTM	GRU	RNN	TCN
Seq. MNIST (accuracy <sup><i>h</i></sup> )	70K	87.2	96.2	21.5	<b>99.0</b>
Permuted MNIST (accuracy)	70K	85.7	87.3	25.3	<b>97.2</b>
Adding problem $T=600$ (loss <sup><math>\ell</math></sup> )	70K	0.164	<b>5.3e-5</b>	0.177	<b>5.8e-5</b>
Copy memory $T=1000$ (loss)	16K	0.0204	0.0197	0.0202	<b>3.5e-5</b>
Music JSB Chorales (loss)	300K	8.45	8.43	8.91	<b>8.10</b>
Music Nottingham (loss)	1M	3.29	3.46	4.05	<b>3.07</b>
Word-level PTB (perplexity <sup><math>\ell</math></sup> )	13M	<b>78.93</b>	92.48	114.50	88.68
Word-level Wiki-103 (perplexity)	-	48.4	-	-	<b>45.19</b>
Word-level LAMBADA (perplexity)	-	4186	-	14725	<b>1279</b>
Char-level PTB (bpc <sup><math>\ell</math></sup> )	3M	1.36	1.37	1.48	<b>1.31</b>
Char-level text8 (bpc)	5M	1.50	1.53	1.69	<b>1.45</b>

Source: [arxiv.org/abs/1803.01271](https://arxiv.org/abs/1803.01271)

# Memory size of TCN and RNNs



*Figure 5.* Accuracy on the copy memory task for sequences of different lengths  $T$ . While TCN exhibits 100% accuracy for all sequence lengths, the LSTM and GRU degenerate to random guessing as  $T$  grows.

Source: [arxiv.org/abs/1803.01271](https://arxiv.org/abs/1803.01271)

# Comparison conclusions

Generic TCNs outperform generic RNNs, LSTMs and GRUs.

Advanced schemes for improving LSTMs have been proposed, while “TCN has not yet benefited from such community-wide investment”.

# Graph2Seq Encoder

# Problems with Seq2Seq

- Sequences may be the simplest structured data
- Some objects are more naturally represented as graphs, which provide more structural information

# Other encoders

- TreeLSTM
- Set2Seq
- Tree2Seq
- Graph2Seq using Gated Graph NNs
- Graph Convolutional Networks

# Architecture example

- ① Graph encoder
  - ① Node embeddings
  - ② Graph embeddings
- ② Sequence decoder takes the embeddings and employs attention over node embeddings while generating sequences

# Node embedding generation

---

**Algorithm 1** Node embedding generation algorithm
 

---

**Input:** Graph  $\mathcal{G}(\mathcal{V}, \mathcal{E})$ ; node initial feature vector  $\mathbf{a}_v, \forall v \in \mathcal{V}$ ; hops  $K$ ; weight matrices  $\mathbf{W}^k, \forall k \in \{1, \dots, K\}$ ; non-linearity  $\sigma$ ; aggregator functions  $\text{AGGREGATE}_k^{\vdash}, \text{AGGREGATE}_k^{\dashv}, \forall k \in \{1, \dots, K\}$ ; neighborhood functions  $\mathcal{N}_{\vdash}, \mathcal{N}_{\dashv}$

**Output:** Vector representations  $z_v$  for all  $v \in \mathcal{V}$

```

1:  $\mathbf{h}_{v\vdash}^0 \leftarrow \mathbf{a}_v, \forall v \in \mathcal{V}$ 
2:  $\mathbf{h}_{v\dashv}^0 \leftarrow \mathbf{a}_v, \forall v \in \mathcal{V}$ 
3: for all  $k = 1 \dots K$  do
4:   for all  $v \in \mathcal{V}$  do
5:      $\mathbf{h}_{\mathcal{N}_{\vdash}(v)}^k \leftarrow \text{AGGREGATE}_k^{\vdash}(\{\mathbf{h}_{u\vdash}^{k-1}, \forall u \in \mathcal{N}_{\vdash}(v)\})$ 
6:      $\mathbf{h}_{v\vdash}^k \leftarrow \sigma(\mathbf{W}^k \cdot \text{CONCAT}(\mathbf{h}_{v\vdash}^{k-1}, \mathbf{h}_{\mathcal{N}_{\vdash}(v)}^k))$ 
7:      $\mathbf{h}_{\mathcal{N}_{\dashv}(v)}^k \leftarrow \text{AGGREGATE}_k^{\dashv}(\{\mathbf{h}_{u\dashv}^{k-1}, \forall u \in \mathcal{N}_{\dashv}(v)\})$ 
8:      $\mathbf{h}_{v\dashv}^k \leftarrow \sigma(\mathbf{W}^k \cdot \text{CONCAT}(\mathbf{h}_{v\dashv}^{k-1}, \mathbf{h}_{\mathcal{N}_{\dashv}(v)}^k))$ 
9:   end for
10: end for
11:  $\mathbf{z}_v \leftarrow \text{CONCAT}(\mathbf{h}_{v\vdash}^K, \mathbf{h}_{v\dashv}^K), \forall v \in \mathcal{V}$ 

```

---

Source: [arxiv.org/abs/1804.00823](https://arxiv.org/abs/1804.00823)



# Aggregate functions

Must be invariant to input permutations.

- Element-wise mean
- LSTM on random permutations of neighbours
- Pooling:  
(neighbour's vector  $\rightarrow$  fully-connected NN)  $\cdot |\mathcal{N}(v)| \rightarrow$  max-pooling

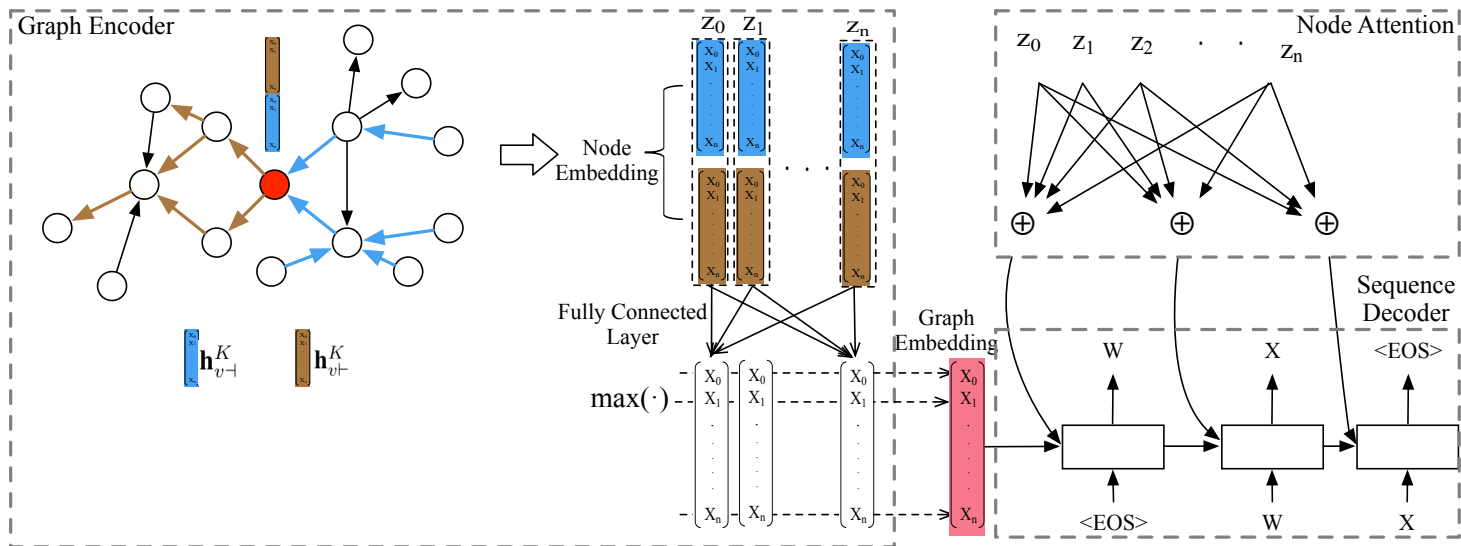
# Graph embedding generation

Graph embeddings are generated from node embeddings.

Two ways:

- ① Pooling-based: max-pooling, min-pooling, average-pooling (feeding them to a fully-connected NN did not reveal any significant performance difference, thus max-pooling is the default method)
- ② Node-based:  
Add a super node,  $v_s$ , and direct the others to it. The node embedding of  $v_s$  obtained by aggregating the embeddings of its neighbour nodes is the target graph embedding

# Scheme



Source: [arxiv.org/abs/1804.00823](https://arxiv.org/abs/1804.00823)

# Settings

- Adam optimiser, mini-batch size = 50
- learning rate = 0.001
- dropout at the decoder level, ratio = 0.5
- clipped gradients, if  $\|\nabla\| > 20$
- hop size  $K = 6$
- $\mathbf{a}_v \in \mathbb{R}^{40}$
- $\sigma = \text{ReLU}$
- decoder: 1 layer, 80 hidden states

Default modification: mean aggregator and pooling-based graph embeddings.

# bAbI Task 19

The bAbI artificial intelligence task 19 (path finding)

garden (A) bathroom (B) bedroom (C)  
hallway (D) office (E) kitchen (F)

1 The **garden** is west of the **bathroom**.  
2 The **bedroom** is north of the **hallway**.  
3 The **office** is south of the **hallway**.  
4 The **bathroom** is north of the **bedroom**.  
5 The **kitchen** is east of the **bedroom**.

Transform

A west B  
B north D  
E south D  
B north C  
F east C

Q: How do you go from the **bathroom** to the **hallway**

Transform

Q:path(B, D)

Source: [arxiv.org/abs/1804.00823](https://arxiv.org/abs/1804.00823)

# Shortest Path Task

Goal: find the shortest directed path between two nodes in a graph.  
SP-S (node size = 5), SP-L (node size = 100).

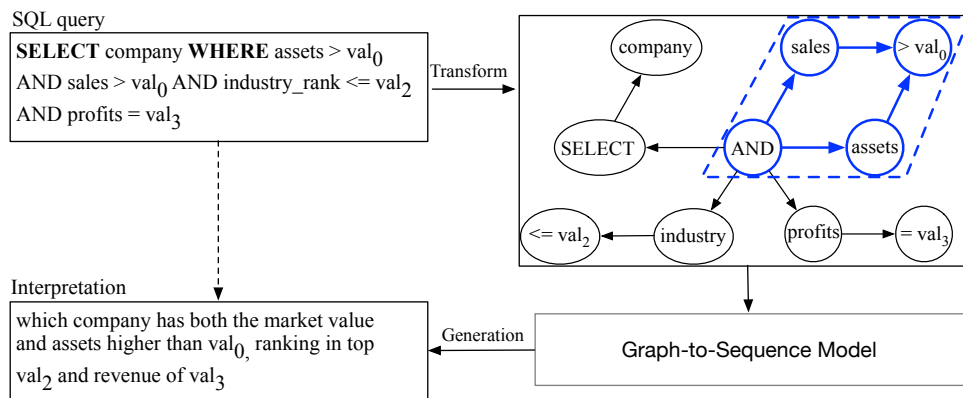
	bAbI T19	SP-S	SP-L
LSTM	25.2%	8.1%	2.2%
GGs-NN	98.1%	100.0%	95.2%
GCN	97.4%	100.0%	96.5%
Graph2Seq	<b>99.9%</b>	100.0%	<b>99.3%</b>

Table 1: Results of our model and baselines on bAbI and Shortest Directed Path tasks.

Source: [arxiv.org/abs/1804.00823](https://arxiv.org/abs/1804.00823)

# Natural language generation

Goal: translate an SQL query to a natural language description expressing its meaning.



Source: [arxiv.org/abs/1804.00823](https://arxiv.org/abs/1804.00823)

	BLEU-4
Seq2Seq	20.91
Seq2Seq + Copy	24.12
Tree2Seq	26.67
GCN-PGE	35.99
GGs-NN	35.53
Graph2Seq-NGE	34.28
Graph2Seq-PGE	<b>38.97</b>

Table 2: Results on WikiSQL.

# ToDo

This version of Graph2Seq encoding needs more testing on other data



# Conclusions

- CNNs work well on NLP tasks
- A single filter often detects numerous semantic classes of ngrams
- Max-pooling induces a threshold behaviour
- Generic TCNs outperform generic RNNs across a broad range of sequence modelling tasks
- Graph2Seq is better suited for encoding certain types of structured data than Seq2Seq

# References

- CNN in NLP in general:
  - General description <http://www.wildml.com/2015/11/understanding-convolutional-neural-networks-for-nlp/>
  - Understanding CNNs in NLP  
<https://arxiv.org/abs/1809.08037>
- CNNs vs RNNs: <https://arxiv.org/abs/1803.01271>
- Graph2Seq: <https://arxiv.org/abs/1804.00823>