


# Improving Stability in Deep Reinforcement Learning with Weight Averaging

Evgenii Nikishin

10.09.2018

# RL problem statement

Markov Decision Process (MDP):



- Environment states  $s_t \in \mathcal{S}$
- Agent actions  $a_t \in \mathcal{A}$
- Reward  $r(s_t, a_t) \in \mathbb{R}$
- Agent policy  $a_t \sim \pi(a_t | s_t)$
- State transitions  $s_{t+1} \sim p(s_{t+1} | s_t, a_t)$

# RL problem statement

Interaction with an environment produces trajectory  $\tau$ :

$$p_{\pi}(\tau) = p(s_0) \prod_{t=0}^T \pi(a_t | s_t) p(s_{t+1} | s_t, a_t)$$

Optimal policy maximizes the expected discounted return:

$$\pi^* = \arg \max_{\pi} \mathbb{E}_{\pi} \left[ \sum_{t=0}^T \gamma^t r(s_t, a_t) \right] = \arg \max_{\pi} \mathbb{E}_{\tau \sim p_{\pi}(\tau)} \left[ \sum_{t=0}^T \gamma^t r(s_t, a_t) \right]$$

$0 \leq \gamma \leq 1$  (0.99 is common)

# Policy gradient methods

Parametrize policy  $\pi(a|s) = \pi_\theta(a|s)$  within a family of differentiable functions and update w.r.t. its gradient:

$$\nabla_{\theta} \mathbb{E}_{\tau \sim p_{\theta}(\tau)} \left[ \sum_{t=0}^T \gamma^t r(s_t, a_t) \right]$$

# Policy gradient methods

Parametrize policy  $\pi(a|s) = \pi_\theta(a|s)$  within a family of differentiable functions and update w.r.t. its gradient:

$$\nabla_\theta \mathbb{E}_{\tau \sim p_\theta(\tau)} \left[ \sum_{t=0}^T \gamma^t r(s_t, a_t) \right] = \sum_{\tau} \nabla_\theta p_\theta(\tau) \sum_{t=0}^T \gamma^t r(s_t, a_t) =$$

## Policy gradient methods

Parametrize policy  $\pi(a|s) = \pi_\theta(a|s)$  within a family of differentiable functions and update w.r.t. its gradient:

$$\begin{aligned}\nabla_\theta \mathbb{E}_{\tau \sim p_\theta(\tau)} \left[ \sum_{t=0}^T \gamma^t r(s_t, a_t) \right] &= \sum_{\tau} \nabla_\theta p_\theta(\tau) \sum_{t=0}^T \gamma^t r(s_t, a_t) = \\ \sum_{\tau} p_\theta(\tau) \nabla_\theta \log p_\theta(\tau) \sum_{t=0}^T \gamma^t r(s_t, a_t) &= \end{aligned}$$

## Policy gradient methods

Parametrize policy  $\pi(a|s) = \pi_\theta(a|s)$  within a family of differentiable functions and update w.r.t. its gradient:

$$\begin{aligned}\nabla_\theta \mathbb{E}_{\tau \sim p_\theta(\tau)} \left[ \sum_{t=0}^T \gamma^t r(s_t, a_t) \right] &= \sum_{\tau} \nabla_\theta p_\theta(\tau) \sum_{t=0}^T \gamma^t r(s_t, a_t) = \\ \sum_{\tau} p_\theta(\tau) \nabla_\theta \log p_\theta(\tau) \sum_{t=0}^T \gamma^t r(s_t, a_t) &= \\ \left\{ \log p_\theta(\tau) = \log p(s_0) + \sum_{t=0}^T (\log \pi_\theta(a_t|s_t) + \log p(s_{t+1}|s_t, a_t)) \right\} &= \end{aligned}$$

## Policy gradient methods

Parametrize policy  $\pi(a|s) = \pi_\theta(a|s)$  within a family of differentiable functions and update w.r.t. its gradient:

$$\begin{aligned}\nabla_\theta \mathbb{E}_{\tau \sim p_\theta(\tau)} \left[ \sum_{t=0}^T \gamma^t r(s_t, a_t) \right] &= \sum_{\tau} \nabla_\theta p_\theta(\tau) \sum_{t=0}^T \gamma^t r(s_t, a_t) = \\ \sum_{\tau} p_\theta(\tau) \nabla_\theta \log p_\theta(\tau) \sum_{t=0}^T \gamma^t r(s_t, a_t) &= \\ \left\{ \log p_\theta(\tau) = \log p(s_0) + \sum_{t=0}^T (\log \pi_\theta(a_t|s_t) + \log p(s_{t+1}|s_t, a_t)) \right\} &= \\ \mathbb{E}_{\tau \sim p_\theta(\tau)} \left[ \left( \sum_{t=0}^T \nabla_\theta \log \pi_\theta(a_t|s_t) \right) \left( \sum_{t=0}^T \gamma^t r(s_t, a_t) \right) \right]\end{aligned}$$



## Policy gradient methods

$$\mathbb{E}_{\tau \sim p_{\theta}(\tau)} \left[ \left( \sum_{t=0}^T \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \right) \left( \sum_{t=0}^T \gamma^t r(s_t, a_t) \right) \right]$$

Main problem: large variance of gradient estimates (much bigger than in supervised learning)

# Policy gradient methods

$$\mathbb{E}_{\tau \sim p_{\theta}(\tau)} \left[ \left( \sum_{t=0}^T \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \right) \left( \sum_{t=0}^T \gamma^t r(s_t, a_t) \right) \right]$$

Main problem: large variance of gradient estimates (much bigger than in supervised learning) Example:

$$r(s_t, a_t) = \begin{cases} -110 & \text{prob } 0.5 \\ -100 & \text{prob } 0.5 \end{cases} \quad \text{Mean} = -105, \text{ Var} = 25$$

$$\nabla_{\theta} \log \pi_{\theta}(a_t | s_t) = \begin{cases} -1, & \text{prob } 0.5 \\ +1, & \text{prob } 0.5 \end{cases} \quad \text{Mean} = 0, \text{ Var} = 1$$

$$\nabla_{\theta} \log \pi_{\theta}(a_t | s_t) r(s_t, a_t) = \begin{cases} -110, & \text{prob } 0.25 \\ -100, & \text{prob } 0.25 \\ +100, & \text{prob } 0.25 \\ +110, & \text{prob } 0.25 \end{cases} \quad \text{Mean} = 0, \text{ Var} = 11050$$

# Dealing with noise

Possible ways to alleviate effect of large variance:

- 1 Introduce a baseline, that does not change the expectation, but can decrease variance (e.g. average reward)

# Dealing with noise

Possible ways to alleviate effect of large variance:

- 1 Introduce a baseline, that does not change the expectation, but can decrease variance (e.g. average reward)
- 2 Introduce some bias in order to decrease the variance (Actor-Critic algorithm [Mnih et al., 2016])

# Dealing with noise

Possible ways to alleviate effect of large variance:

- ➊ Introduce a baseline, that does not change the expectation, but can decrease variance (e.g. average reward)
- ➋ Introduce some bias in order to decrease the variance (Actor-Critic algorithm [Mnih et al., 2016])
- ➌ Penalise for large policy deviations (Trust Region Policy Optimization [Schulman et al., 2015])

# Dealing with noise

Possible ways to alleviate effect of large variance:

- ➊ Introduce a baseline, that does not change the expectation, but can decrease variance (e.g. average reward)
- ➋ Introduce some bias in order to decrease the variance (Actor-Critic algorithm [Mnih et al., 2016])
- ➌ Penalise for large policy deviations (Trust Region Policy Optimization [Schulman et al., 2015])
- ➍ ...more?

Stochastic Weight Averaging [Izmailov et al., 2018]: average the weights collected during training with an SGD-like method.


- 1 start after conventional pretraining
- 2 use constant or cyclical learning rate  
(to continue exploring regions of high-performing networks)
- 3 dynamically recalculate average:

$$w_{\text{SWA}} \leftarrow \frac{n_{\text{SWA}} \cdot w_{\text{SWA}} + w}{n_{\text{SWA}} + 1}$$

$$n_{\text{SWA}} \leftarrow n_{\text{SWA}} + 1$$

# Motivation in SL



Mandt et al. [2017]: SGD with fixed learning rate samples from a Gaussian distribution centered at the minimum of the loss, i.e. SGD iterates stay at the boundary of a high-quality region:





# Motivation in SL


Due to shift between train and test loss surfaces, we are looking for wider optima. SWA leads to better generalization:



# Motivation in RL

Deep RL methods are notoriously unstable.

A2C cumulative reward trajectory on CartPole environment with and without weight averaging:



SWA alleviates the effect of noise during training.

# SWA for A2C and DDPG

We apply SWA to

- Advantage Actor-Critic [Mnih et al., 2016] for discrete action space environments (Atari games).
- Deep Deterministic Policy Gradient [Lillicrap et al., 2015] for continuous action space environments (MuJoCo).

After pretraining with conventional training, we apply SWA by collecting weights every  $c$  timesteps.

# Atari experiments

Table: Average final cumulative reward with and without SWA.

ENV NAME	A2C	A2C + SWA
Breakout	522 $\pm$ 34	<b>703 <math>\pm</math> 60</b>
Qbert	18777 $\pm$ 778	<b>21272 <math>\pm</math> 655</b>
SpaceInvaders	7727 $\pm$ 1121	<b>21676 <math>\pm</math> 8897</b>
Seaquest	1779 $\pm$ 4	<b>1795 <math>\pm</math> 4</b>
CrazyClimber	<b>147030 <math>\pm</math> 10239</b>	139752 $\pm$ 11618
BeamRider	9999 $\pm$ 402	<b>11321 <math>\pm</math> 1065</b>

Table: Average final cumulative reward with and without SWA.

ENV NAME	DDPG	DDPG + SWA
Hopper	$613 \pm 683$	<b><math>1615 \pm 1143</math></b>
Walker2d	$1803 \pm 96$	<b><math>2457 \pm 241</math></b>
Half-Cheetah	$3825 \pm 1187$	<b><math>4228 \pm 1117</math></b>
Ant	$865 \pm 899$	<b><math>1051 \pm 696</math></b>

RL + SWA:

- Easy to implement
- Sample-efficient way to improve stability
- Alleviates problem of forgetting good policies

# References

- Pavel Izmailov, Dmitrii Podoprikin, Timur Garipov, Dmitry Vetrov, and Andrew Gordon Wilson. Averaging weights leads to wider optima and better generalization. *arXiv preprint arXiv:1803.05407*, 2018.
- Timothy P Lillicrap, Jonathan J Hunt, Alexander Pritzel, Nicolas Heess, Tom Erez, Yuval Tassa, David Silver, and Daan Wierstra. Continuous control with deep reinforcement learning. *arXiv preprint arXiv:1509.02971*, 2015.
- Stephan Mandt, Matthew D Hoffman, and David M Blei. Stochastic gradient descent as approximate bayesian inference. *The Journal of Machine Learning Research*, 18(1):4873–4907, 2017.
- Volodymyr Mnih, Adria Puigdomenech Badia, Mehdi Mirza, Alex Graves, Timothy Lillicrap, Tim Harley, David Silver, and Koray Kavukcuoglu. Asynchronous methods for deep reinforcement learning. In *International Conference on Machine Learning*, pages 1928–1937, 2016.
- John Schulman, Sergey Levine, Pieter Abbeel, Michael Jordan, and Philipp Moritz. Trust region policy optimization. In *International Conference on Machine Learning*, pages 1889–1897, 2015.